Data Science and the Data Scientist Toolkit

// FLATIRON SCHOOL

Agenda

- What is Data Science?
  - Roles and Responsibilities
  - The Process
- The Data Science Toolkit (Phase 1)

So:
What is
Data Science?

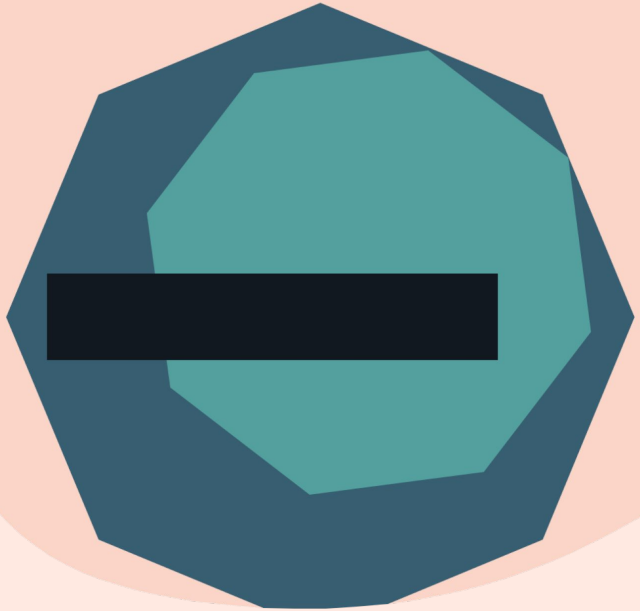

What is Data Science?

Find out for yourself!

Prompt: Spend the next 10 minutes skimming and discussing your assigned blog post, then come back and report your findings to the rest of us.

1. A Deep Look Into 13 Data Scientist Roles and Their Responsibilities
2. The Data Science Process
3. Most In Demand Data Science Technical Skills
4. A Learning Path to Becoming a Data Scientist
5. Compilation of Advice for New and Aspiring Data Scientists

Let's Discuss!

What does a "data scientist" do?

What are the main skills you need to be a "data scientist" ?

What is consistent among these posts, and what is in dispute?

Let's Discuss!

1: Different roles in DS, depends on individual and technical ability

-Tech heavy roles

-Business heavy roles

2: Acting as liaison b/t tech and product (sales)

-Customer retention (churn)

-Summarize technical into nontechnical

3: Manual labor that goes into data automation

-Data cleaning

-Optimization

-Interaction of lots of moving parts

4: Learning path to DS

-Continuous learning process

-Version control, package mgm

-All the way to ML

5: Key themes to keep in mind

-Communication

-Statistics

-Question and confirm

-make sure work public

“A data scientist is a professional responsible for **collecting, analyzing and interpreting** extremely large amounts of data. The data scientist role is an **offshoot of several traditional technical roles**, including mathematician, scientist, statistician and computer professional.”

The Data Science Venn Diagram

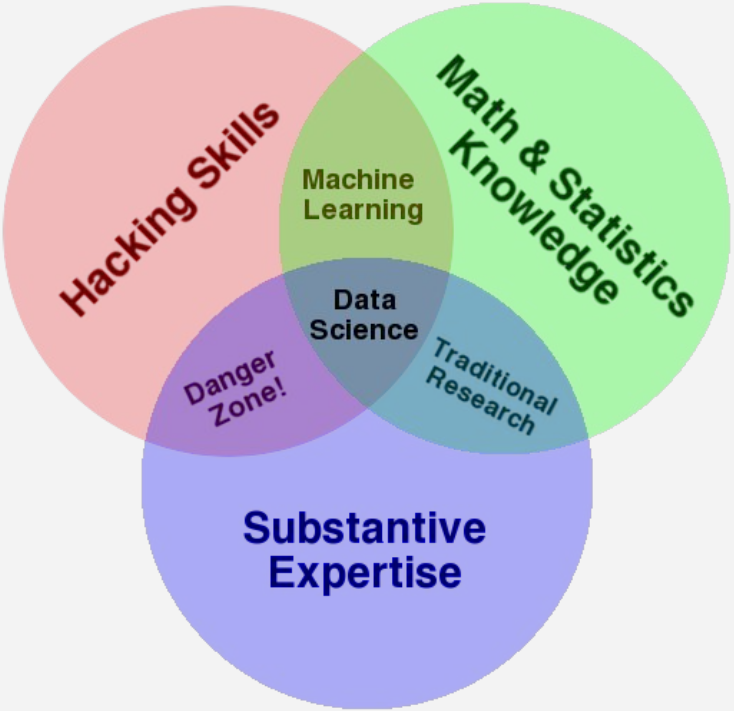

Another Version

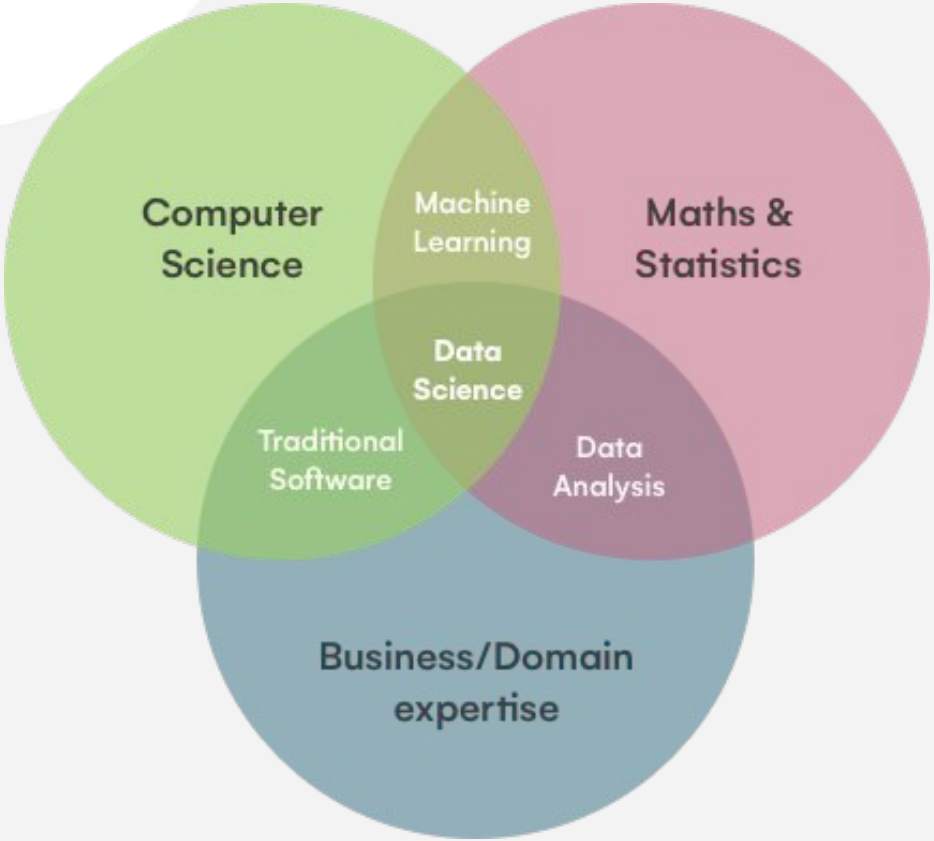

And Another

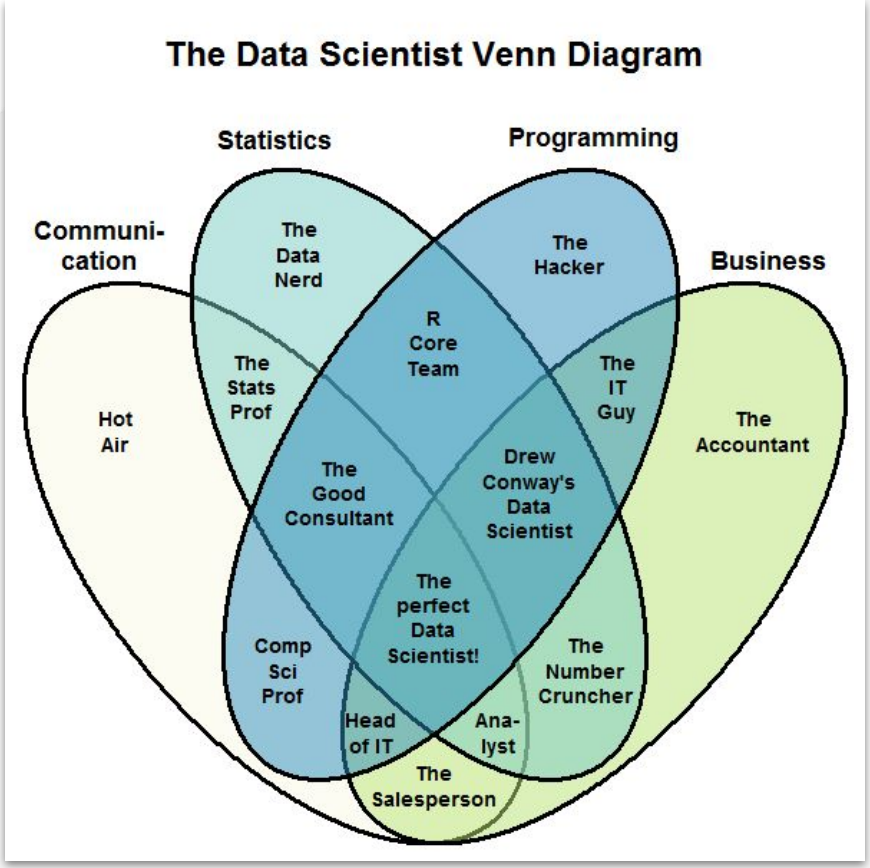
The Data Scientist Venn Diagram
Statistics
Programming
Communi-cation
Business
The Data Nerd
The Hacker
R Core Team
The Stats Prof
The IT Guy
Hot Air
The Accountant
The Good Consultant
Drew Conway's Data Scientist
The perfect Data Scientist!
Comp Sci Prof
The Number Cruncher
Head of IT
Ana-lyst
The Salesperson

Common Roles & Responsibilities

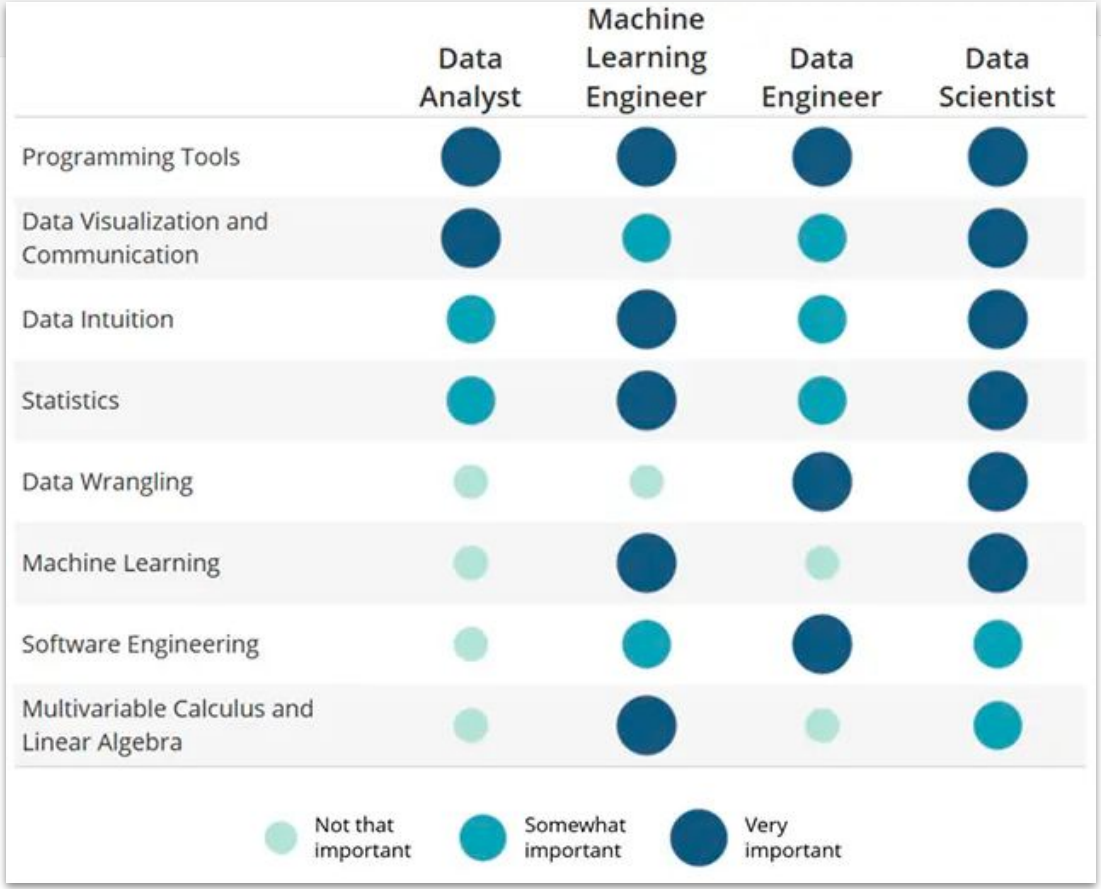

“Regardless of your exact job title, if you’re in the field of data science, you’ll be expected to be **involved in a lot of different steps** in the data-driven product development cycle. You should be ready to discover new areas to **optimize**, figure out the **metrics** that matter, find the **data** to inform these metrics, design and execute **experiments**, and **present the results** of experiments/models in concise, accurate, and convincing ways.”

The Data Science Process

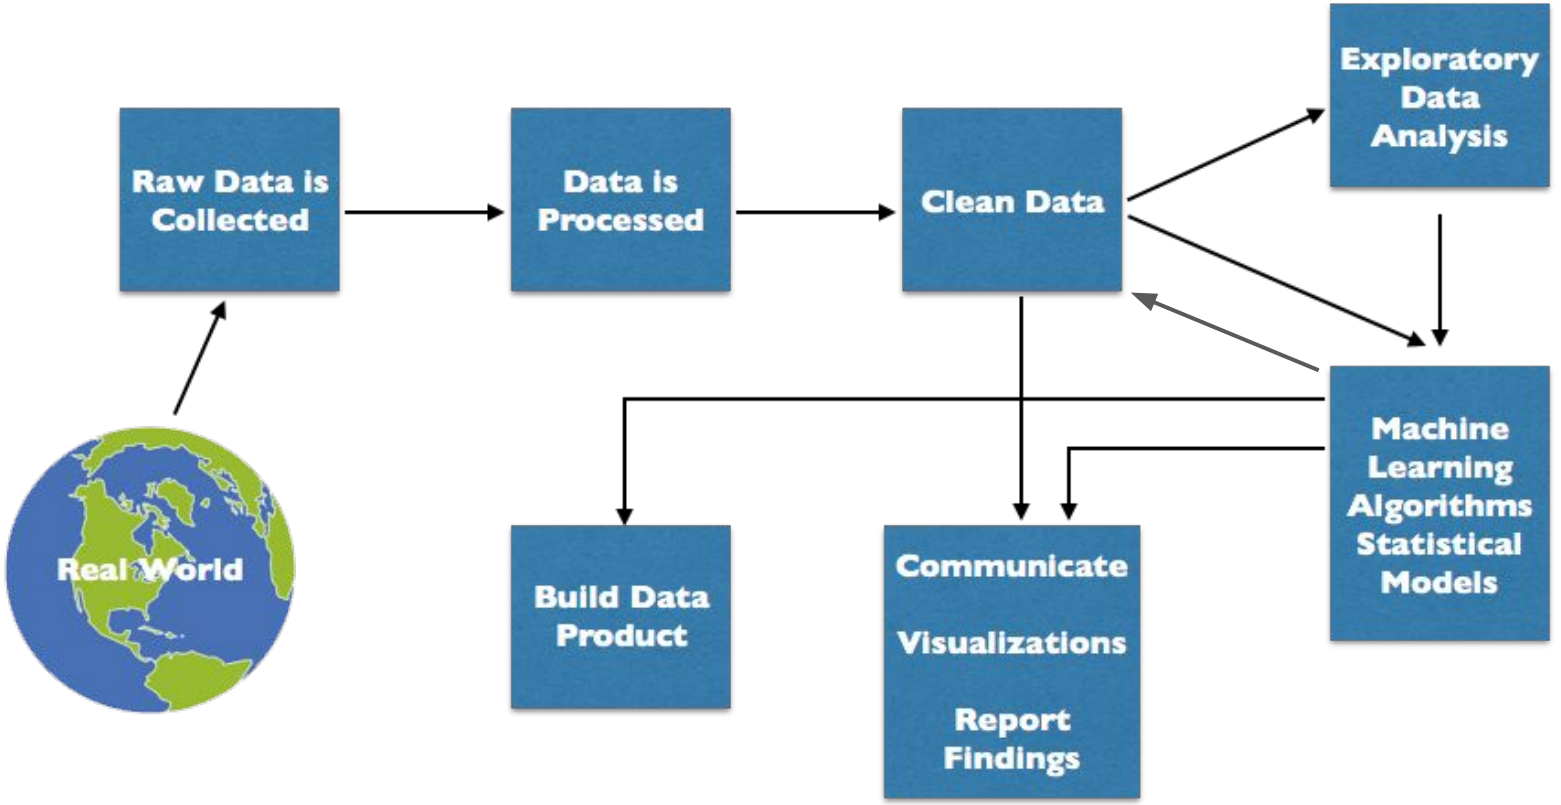

And Another

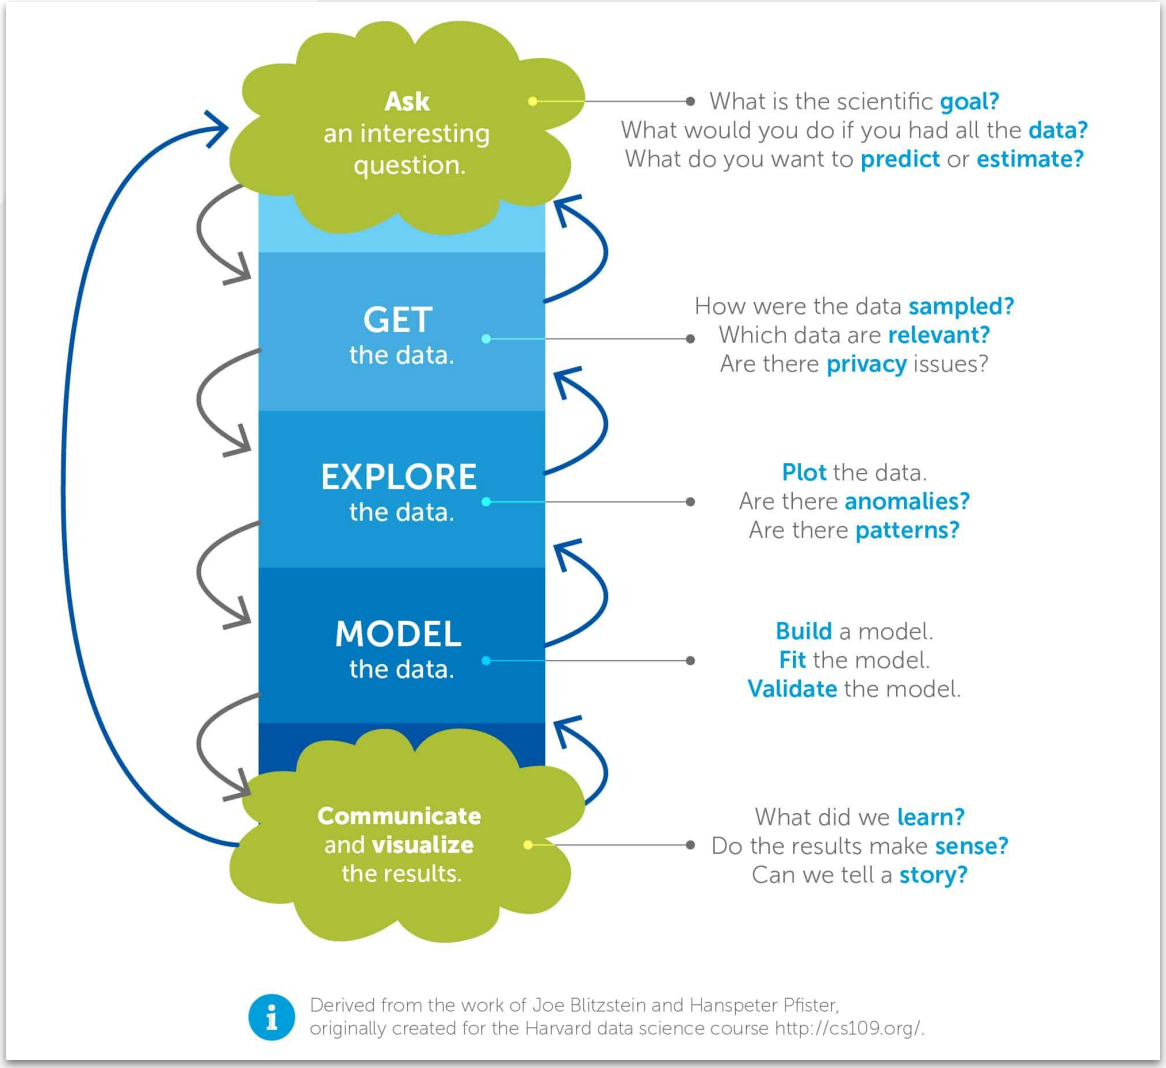
Ask
an interesting
question.
What is the scientific goal?
What would you do if you had all the data?
What do you want to predict or estimate?
GET
the data.
How were the data sampled?
Which data are relevant?
Are there privacy issues?
EXPLORE
the data.
Plot the data.
Are there anomalies?
Are there patterns?
MODEL
the data.
Build a model.
Fit the model.
Validate the model.
Communicate
and visualize
the results.
What did we learn?
Do the results make sense?
Can we tell a story?
Derived from the work of Joe Blitzstein and Hanspeter Pfister,
originally created for the Harvard data science course http://cs109.org/.

CRISP-DM Process Diagram

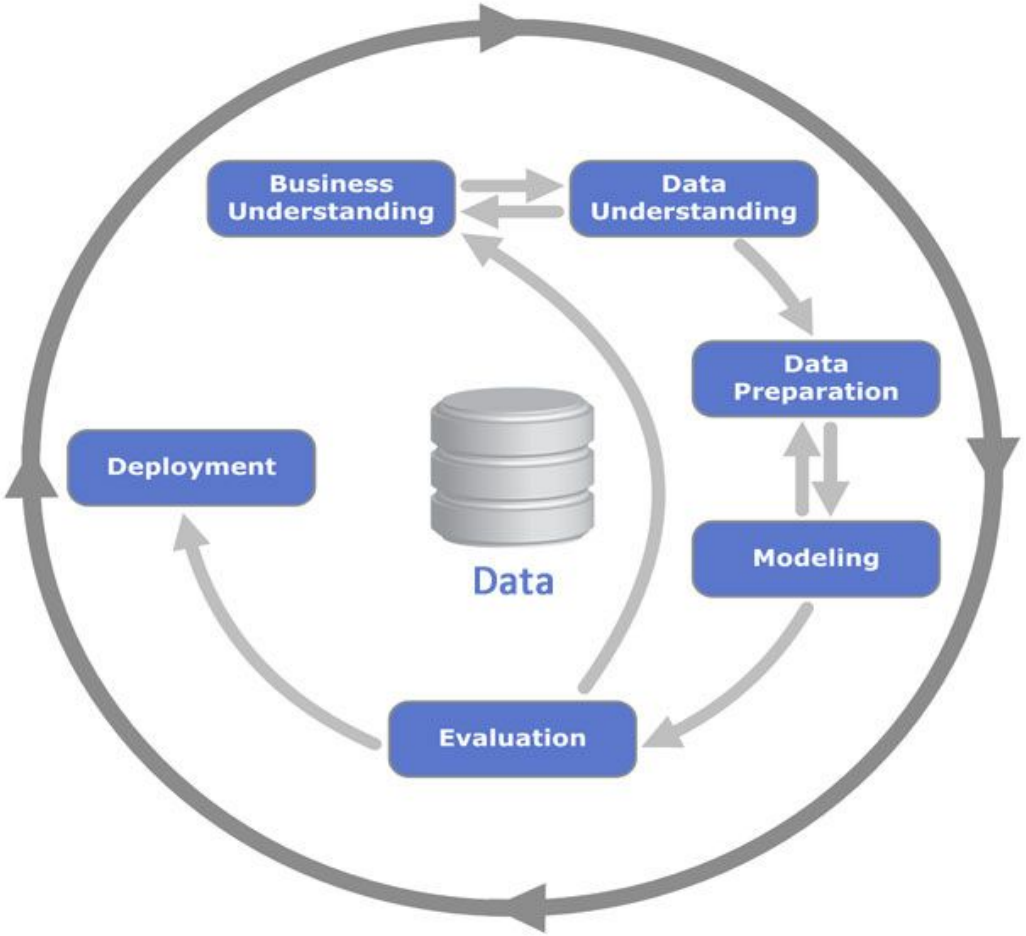

Source: Kenneth Jensen

Data Science Process

LEAD

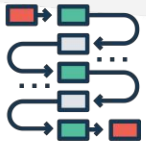

OBTAIN	SCRUB	EXPLORE	MODEL	INTERPRET
O	**S**	**E**	**M**	**N**
Gather data from relevant sources	Clean data to formats that machine understands	Find significant patterns and trends using statistical methods	Construct models to predict and forecast	Put the results into good use

Originally by Hilary Mason and Chris Wiggins

The Data Science Toolkit

Data Science Toolkit - Phase 1

Languages

Interfaces

Version Control

Package Control

Languages

Python

- Free, open source, versatile, powerful
- Not just for data science!
- Object-oriented (everything is an ‘object’)
- The Zen of Python

Structured Query Language (SQL)

- Connect to, change, and retrieve data from relational databases
- Developed in the 1970s, still going strong
- Many flavors

Interfaces

Jupyter Notebooks

- Streamlined document-centric interface for running and sharing code

IllumiDesk

- Hosts Jupyter Notebooks in the cloud

Code-Focused Text Editor

- Write text files in a code-native format
- **VS Code** is one of many that would work

Version Control

Git

- Distributed version tracking on any files
- Folder → “Repository”

GitHub

- Hosts Git repositories
- Collaborate and share code with others
- Backbone of the open source community
- Your Data Science portfolio!

Package Control

Anaconda

- Package management and deployment
- Designed with Data Science in mind
- Create and share environments

Python Package Index (PyPi)

- Database of public Python libraries
- Package installer (pip)
- Not everything is on Anaconda

Now:
Time to Get Started!